

Distance sampling online workshop

Analysis in R: Exploratory analysis

CREEM, Univ of St Andrews – October 2018

1 R basics

1.1 Introduction

This tutorial is the first in a series that were created as an introduction for those wishing to undertake distance sampling and related methods in R ([R Core Team, 2018](#)). The basics of R and RStudio are explored in this tutorial. Data are generated resembling a set of perpendicular distances collected from a track line and basic R commands are used to summarise and analysis it.

This practical is divided into three; the first part is a basic introduction to R and RStudio, the second part is an introduction to R commands and the final part is an introduction to the distance sampling package.

Throughout these practicals, R commands are shown in `fixed pitch` font and text in *italics* which starts with a hash is a comment, for example

```
# This is an example of a comment
```

1.2 Objectives of this practical

1. Introduction to R and RStudio.
2. Using basic functions to summarise data.
3. Displaying data using histograms.
4. Performing basic calculations.
5. Using the distance sampling package to fit a detection function model.

1.3 Introduction to RStudio

RStudio is a graphical user interface (GUI) to R; a major advantage of using RStudio is that it makes it easy to type R code into a script window which can be sent to the R command line to be executed and can also be saved (and hence retrieved). This section describes how to open and use the user-interface RStudio.

Check that RStudio and R have been installed on your computer; for this course we are using R version 3.5.1. When you open an RStudio session, the window will likely be similar to that shown in Figure 1.

RStudio has four main panel sections;

- a script panel (top left),
- R-console (bottom left),
- Environment and History panel (top right) and
- a panel for Files, Plots, Packages and Help (bottom right).

If you can only see three panels, go to 'File > New File > R Script' to open a script panel.



"Bigeye binoculars used in a cetacean survey"

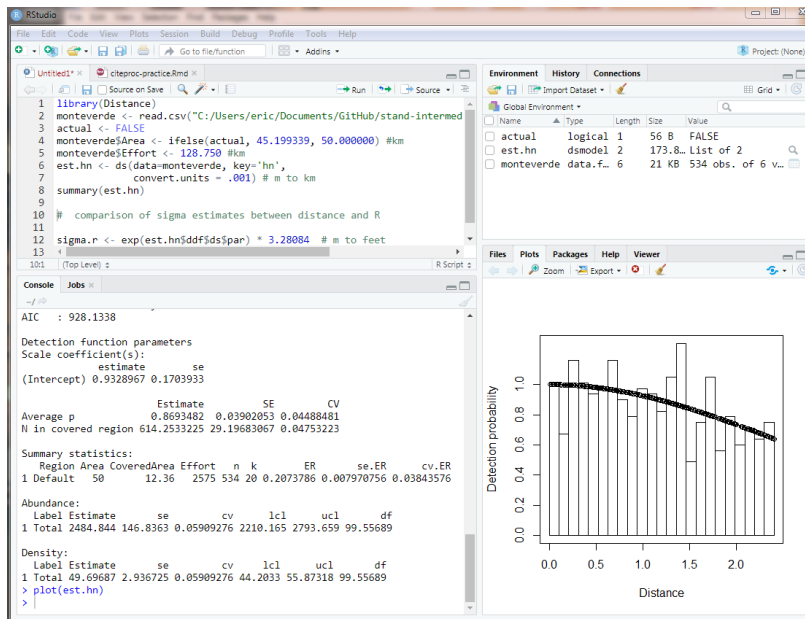


Figure 1: R-Studio screenshot

The script panel is used for writing R commands and comments describing the code to help you remember what the code does and why. The script can be saved and revisited at a later date; this is useful if writing a lot of code. The saved script file has a suffix 'xxx.Rmd', where xxx is a name of your choosing.

Directly below the script panel is the R console window. R is the program that performs the analysis and executes your calculations. There are several options for executing code:

- to run code from your script window, move the cursor to the line of code to be run and click the 'Run' button in the top right of the script window,
- or use keyboard shortcuts for Windows = Ctrl+Enter, for Mac/Linux = command+enter,
- copy the commands from the script panel and paste them into the R console panel at the > symbol,
- type commands directly into the console panel at the > symbol.

The top right panel contains two tabs, Environment and History. The Environment tab shows all the objects and functions that have been imported and created in the R workspace. Upon exiting RStudio, the R workspace can be saved and then retrieved at a later date. An R object could be data that has been imported into R or objects created within R during an analysis. Also, if you create a function (a collection of commands) during your work session, these will also be listed on the Environment tab. The History tab shows all of the R code that has been run during the session; this is useful if you want to repeat or check a command.

In the bottom right panel there is the Packages tab; R packages are collections of functions which provide tools required to perform particular tasks. For the majority of the practicals for this course, additional packages will not have to be loaded. The Plots tab displays plots produced during the 'R' session. There is also a tab displaying details associated with in-built R func-

tions; the help tab is very useful (see below). Finally, there is a Files tab; this tab lists files available in a particular (i.e. working) directory. The user can point R to a particular directory; this is called the working directory (more on this later).

1.4 The help function

There is a help facility in R which provides details about functions and packages. For example, to learn about the function `round` and what arguments it requires, type the following at the `>` in the console tab:

```
help(round)
```

or alternatively,

```
'?'(round)
```

The description of the function will appear on the Help tab (bottom right). Help is available for any function (even those you write yourself, if you create proper documentation).

1.5 Generating data

To get started, we will generate some data, pretending it is a set of perpendicular distances, and analyse it using basic commands. Below is a command to generate some data - this is fairly complicated but, for now, just copy and paste, or type, the command into the console window.

```
dists <- rnorm(n = 50, mean = 0, sd = 3)
```

At the moment, the syntax of this command can be ignored and will be explained as you work through the practical, nevertheless, what has this command done? It has:

- generated 50 random values from a Normal distribution (using the `rnorm` function) that has a mean of 0 and standard deviation of 3 (i.e. $N(\mu = 0, \sigma^2 = 3^2)$)
- stored the values in an object called `dists`.

The reason to store the values is so we can use them during this practical. We can think of these data as perpendicular distances from a transect centre line; since the mean of the numbers is 0, then some values will be negative (distances on the left of the centre line) and some will be positive (distances on the right). The numbers are generated randomly and so each time this command is executed, the numbers will be different; therefore, the set of values you generate will not be exactly the same as shown here.

Having generated some data, the first thing to do is summarise and plot it.

1.6 Summarising data

Typing the name of an object (in this case `dists`) will list the object on the screen i.e.

```
dists
```

```
## [1]  5.11470965 -2.13611569 -0.83395474
## [4] -0.35894707 -0.37188185  0.80455131
## [7]  2.18052447  0.69940624  1.01734163
## [10] -1.65574400  1.04310407  4.45377542
## [13]  0.56497656  7.32977947 -3.46031846
## [16] -2.41401505  1.36820744  1.26099771
## [19]  1.73275349  1.33906818  2.75176659
## [22] -0.32118462  2.96500619 -3.21671641
## [25] -2.27404584  0.28500216 -6.99279352
## [28]  1.25254794 -3.36098227 -1.42405409
## [31] -4.59123651  1.24714446  1.90259382
## [34]  3.69254206 -0.46369099  0.34361834
## [37] -6.67878992  4.90737102 -0.47796215
## [40]  0.08480238 -4.60995986  0.23447337
## [43]  0.75118743  0.73226160  2.39921469
## [46]  1.02422867  0.77615303  2.86163585
## [49]  3.40924416 -2.95286406
```

Listing all the data is not very helpful, especially if there is a lot of data, but there are ways to display a subset of it.

Display the first six values:

```
head(dists)
```

```
## [1]  5.1147097 -2.1361157 -0.8339547
## [4] -0.3589471 -0.3718819  0.8045513
```

Another way to display the first six values is to select which values to display, for example in the command below the second to eighth value, inclusive, will be displayed,

```
dists[2:8]
```

```
## [1] -2.1361157 -0.8339547 -0.3589471
## [4] -0.3718819  0.8045513  2.1805245
## [7]  0.6994062
```

Listing all (or a few) of the numbers is helpful to see what it looks like but is not very helpful to get an overall picture and so we can use in-built R functions to help summarise it. Mostly, the function names are fairly intuitive:

```
mean(dists)
```

```
## [1] 0.2386946
```

The mean should be zero, or very close to zero, since we specified a mean of zero to generate the data. Other useful functions are given below.

```
min(dists)
```

```
## [1] -6.992794
```

```
max(dists)
```

```
## [1] 7.329779
```

All the above information, and more, is provided in one command as follows:

```
summary(dists)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.
## -6.9928 -1.2765   0.7158   0.2387  1.6416
##      Max.
##      7.3298
```

1.7 Histograms

A key component of distance sampling is the histogram and we can plot our data using the `hist` function:

```
hist(dists)
```

Functions often have arguments which are specified by default and hence do not need to be specified by the user, unless the user wishes to change the default values. For example, the number of intervals in a histogram can be changed using the `breaks` argument (see command below).

Note if we provide a number of intervals for the `breaks` argument, it doesn't always give the specified number of intervals - it works out useful/pretty intervals. Later in the practical we will specify the bin intervals.

We can add a vertical (`v=`), red line to the histogram to indicate the transect centre line:


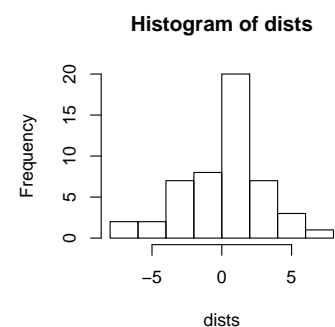
```
hist(dists, breaks = 15)
abline(v = 0, col = "red")
```

In distance sampling, the left and right of the centre line are not differentiated and so to continue we want to ignore the sign of the values and convert negative values to positive values. The command below does this:

```
posdists <- abs(dists)
```

The command can be broken down as follows

- `abs` is a function which takes the absolute value of an object, in this case, it takes the absolute values of `dists`
- `<-` is the assignment symbol which assigns values on the right hand side (i.e. absolute values of `dists`) to the object named on the left hand side (i.e. `posdists`). Plot a histogram of the new object and also specify a label for the *x*-axis using the `xlab` argument:




```
hist(posdists, xlab = "Distance")
```

To specify the bin intervals, first we create the required intervals by generating a sequence of numbers (i.e. 0, 0.5, 1, ..., 8) and then using them in the `hist` command:

```
setbreaks <- seq(from = 0, to = 8, by = 0.5)
hist(posdists, breaks = setbreaks, xlab = "Distance")
```

Histogram of posdists



1.8 Using Distance to fit a detection function model

We can use the `Distance` package (Miller, 2017) to fit a detection function and estimate the probability of detecting objects given they are within the truncation distance of the transect.

There are hundreds of R packages and so for efficiency the more specialised packages are not installed as standard when R is downloaded onto your computer. To check whether the `Distance` package is installed, click on the 'Packages' tab (lower right pane) and see if it is listed. If so, then either click in the tick box or load it using the command

```
library(Distance)
```

If it is not listed, then it can easily be downloaded using the 'Install' tool. Note that each time you start a new R session and wish to use `Distance`, you will need to load it.

The function `ds` fits a detection function and calculates abundance from both line or point transects. In this example, we don't have any areas required to calculate density and abundance and just want to estimate the probability of detection. The distances were generated from a Normal distribution and so a half-normal key function is specified (with the argument `(key="hn")`) with no adjustment terms (`adjustment=NULL`).

```
ds(posdists, key = "hn", adjustment = NULL)
```

If we want to save the output so that we can interrogate it later, then we can assign the output to an object (`dsmodel`) as follows (the `quiet=TRUE` suppresses any message being printed):

```
dsmodel <- ds(posdists, key = "hn", adjustment = NULL,
              quiet = TRUE)
```


To examine the detection function, use the `summary` function. The `summary` function is a generic function and selects a relevant 'summary' based on the type of the object, in this case a distance sampling model object.

```
summary(dsmodel)
```

```
##
```

```
## Summary for distance analysis
```

Histogram of posdists



```
## Number of observations : 50
## Distance range       : 0 - 7.329779
##
## Model : Half-normal key function
## AIC   : 178.0555
##
## Detection function parameters
## Scale coefficient(s):
##           estimate      se
## (Intercept) 1.099874 0.1028484
##
##           Estimate      SE
## Average p      0.5060759 0.04681069
## N in covered region 98.7994057 13.41426836
##
##           CV
## Average p      0.09249736
## N in covered region 0.13577276
```

The function plot is a generic function which can be used to plot the detection function:

```
plot(dsmodel)
```

To truncate the distances, the truncation argument in ds can be used:

```
dstrunc <- ds(posdist, key = "hn", adjustment = NULL,
             truncation = 6)
```

```
summary(dstrunc)
```

```
##
## Summary for distance analysis
## Number of observations : 47
## Distance range       : 0 - 6
##
## Model : Half-normal key function
## AIC   : 149.1578
##
## Detection function parameters
## Scale coefficient(s):
##           estimate      se
## (Intercept) 0.9077179 0.1281165
##
##           Estimate      SE
## Average p      0.5097353 0.05846314
## N in covered region 92.2047200 14.16045945
##
##           CV
## Average p      0.1146931
## N in covered region 0.1535763
```




Figure 2: Half normal detection function fitted to all simulated data.




Figure 3: Half normal detection function fitted to simulated data truncated at distance=6.

References

- Miller, D. L. 2017. Distance: Distance Sampling Detection Function and Abundance Estimation. URL <https://CRAN.R-project.org/package=Distance>. R package version 0.9.7.
- R Core Team. 2018. R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.